

# **Adaptive Mesh Refinement for Solving the Advection-Diffusion Equation**

**A High-Performance Computing Approach**

Kylie Hefner

CS 616: High Performance Computing

December, 2024

# Purpose of the Project

## Main Goal

To develop an efficient and accurate solver for the **advection-diffusion equation** using **adaptive mesh refinement (AMR)** and **parallelization** techniques.



## Key Objectives

- **Numerical Accuracy:** Improve solution precision for dynamic systems.
- **Computational Efficiency:** Reduce unnecessary computations using AMR.
- **Scalability:** Leverage OpenMP and MPI to handle large-scale problems effectively.

# Situation and Challenges

## Situation

The advection-diffusion equation models physical phenomena such as:

- Heat transfer
- Pollution dispersion
- Population dynamics in ecosystems

## 01

### Computational Cost

Uniform grids waste resources in areas without high resolution needs

## 02

### Numerical Instabilities

Advection and diffusion can cause unphysical oscillations

## 03

### Large-Scale Problems


Real-world simulations demand scalability and efficiency

## Problems

# The Advection-Diffusion Equation


## Advection

Describes the transport of a substance (e.g., heat, pollutants, or particles) due to a velocity field



## Diffusion

Describes the spreading of a substance due to random motion, such as molecular collisions



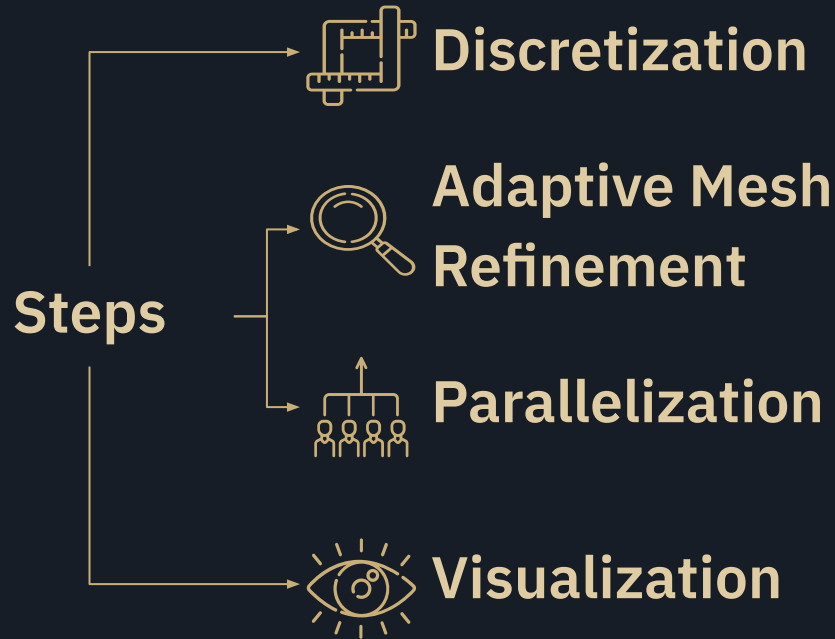
# The Advection-Diffusion Equation

The diagram illustrates the Advection-Diffusion Equation with the following components and annotations:

$$\frac{\partial u}{\partial t} = D \nabla^2 u - \vec{v} \cdot \nabla u$$

- Velocity Vector**: An annotation with a curved arrow pointing to the vector  $\vec{v}$  in the advection term.
- Diffusion Coefficient**: An annotation with a straight arrow pointing to the coefficient  $D$  in the diffusion term.
- Scalar Field (e.g. concentration or temperature)**: An annotation with two arrows pointing to the scalar field  $u$  in both the diffusion and advection terms.

# Methodology




Apply upwind differencing for advection and central differencing for diffusion in C++

Refine grids dynamically in regions with high solution gradients

Use OpenMP and MPI to optimize performance and scalability

Export results and generate heatmaps, contour plots, and animations in Python

# Discretized Equation

$$u_{i,j}^{n+1} = u_{i,j}^n - \Delta t \cdot \left( v_x \cdot \frac{\partial u}{\partial x} + v_y \cdot \frac{\partial u}{\partial y} \right) + \Delta t \cdot D \cdot \frac{\partial^2 u}{\partial x^2}$$
Three orange arrows originate from the text below and point upwards towards the partial derivative terms in the equation. The first arrow points to the  $\frac{\partial u}{\partial x}$  term, the second arrow points to the  $\frac{\partial u}{\partial y}$  term, and the third arrow points to the  $\frac{\partial^2 u}{\partial x^2}$  term.

To approximate partial derivatives, I used  
upwind and central differencing

# Numerical Methods

## Upwind Differencing for Advection

Uses flow direction when calculating spatial derivatives

Prevents instability common with central differencing

$$\frac{\partial u}{\partial x} \approx \frac{u_i - u_{i-1}}{\Delta x}, \quad \text{if } v > 0$$

$$\frac{\partial u}{\partial x} \approx \frac{u_{i+1} - u_i}{\Delta x}, \quad \text{if } v < 0$$

## Central Differencing for Diffusion

Approximates second derivatives symmetrically

Ensures numerical accuracy for diffusion process

$$\frac{\partial^2 u}{\partial x^2} \approx \frac{u_{i+1} - 2u_i + u_{i-1}}{\Delta x^2}$$

## Time Integration

Explicit time-stepping with adaptive  $\Delta t$

$\Delta t$  is dynamically chosen to satisfy stability criteria



# Stability Criteria

## Advection Stability

(Courant-Friedrichs-Lewy condition)

$$\Delta t < \frac{\Delta x}{\max(|v_x|, |v_y|)}$$

## Diffusion Stability

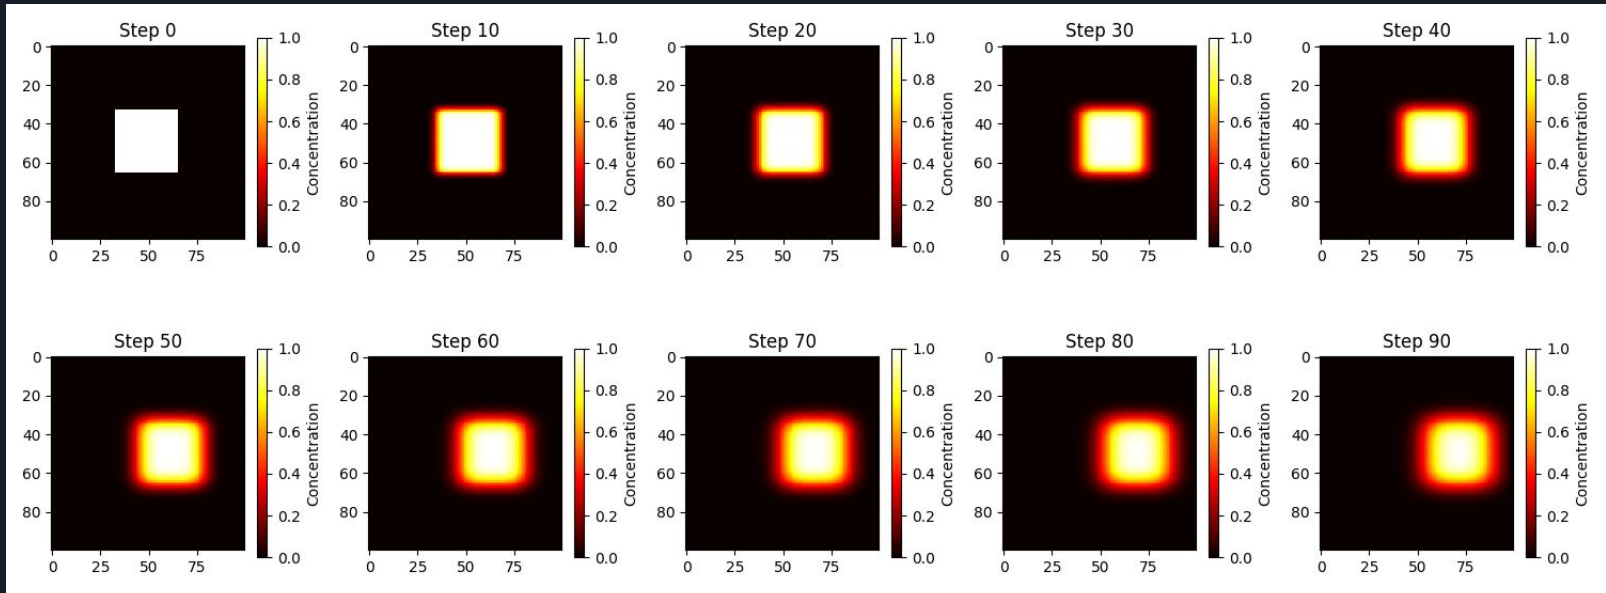
$$\Delta t < \frac{\Delta x^2}{4D}$$

## Combined Stability

$$\Delta t < \min\left(\frac{\Delta x}{\max(|v_x|, |v_y|)}, \frac{\Delta x^2}{4D}\right)$$


# Results

100x100 uniform grid, with heatmap showing the evolution over time.




# Results

100x100 uniform grid, with contour plot showing the gradient over time.



# What if We Saved Computational Power with an Adaptive Grid?




# Quadtree Grid for Adaptive Refinement

What is a Quadtree?

- A hierarchical structure that dynamically divides the domain into smaller cells
- Each parent cell can split into four child cells for higher resolution

Why use a Quadtree?

- Reduces computational cost by refining only where needed
- Ensures accuracy in regions with high gradients



# Error Estimation Based on Gradients

**Objective:** Identify regions requiring refinement based on gradients

**Error Estimator:**

$$\text{Error} = \left| \frac{\partial u}{\partial x} \right| + \left| \frac{\partial u}{\partial y} \right|$$

**Gradient Calculation (Central Differencing):**

$$\frac{\partial u}{\partial x} \approx \frac{u_{i+1,j} - u_{i-1,j}}{2\Delta x}, \quad \frac{\partial u}{\partial y} \approx \frac{u_{i,j+1} - u_{i,j-1}}{2\Delta y}$$

# AMR Implementation

1. Calculate Errors:
  - a. Use gradients to compute the error for each cell.
2. Refine or Coarsen:
  - a. When the error exceeds the threshold, refine cells by dividing a cell into four child cells
  - b. When the error is below the lower threshold coarsen cells by merging child cells back into parent cell
3. Update Solution:
  - a. Apply numerical methods to the refined grid

# Uniform Grid Vs. Adaptive Grid

Uniform Runtime:  
1.95264 s

Adaptive Runtime:  
5.85792 s

Initially, my adaptive grid solver was significantly slower than my uniform grid solver due to excessive overhead. (200x200 starting grid)

Implemented Fixes:

- Apply refinement and coarsening less often (10 time steps instead of every step)
- Increase the refine and coarsen thresholds to reduce computational cost



# Uniform Grid Vs. Adaptive Grid


Uniform Runtime:  
1.95264 s

Adaptive Runtime:  
2.23960 s

After adjusting the refinement and coarsening, the adaptive grid now performs similar to the uniform grid given the same starting size.

**Important Note:** The runtime with AMR is not always faster when using the same initial grid size as a uniform grid. Instead, the goal of AMR is to focus computational effort on areas of interest (regions with high gradients or errors) while maintaining lower resolution elsewhere.

# Uniform Grid Vs. Adaptive Grid



## Example Use Case:

- A uniform grid of  $64 \times 64$  might need to be fully refined to capture small-scale features, leading to high runtime and memory usage.
- With AMR, you can start with a  $8 \times 8$  grid and refine locally, achieving similar accuracy with fewer cells.

# Parallelization with OpenMP

## Parallelizing Solution Updates:

- Each advection-diffusion grid cell is updated independently, so we can use openmp threads to parallelize the nested loops

#pragma omp parallel for collapse(2)

- Parallelizes the two nested loops over i and j
- collapse(2) combines the two loops into a single iteration space for efficient scheduling

```
// Function to update the grid values based on advection-diffusion equation
void update_solution(Matrix &grid, const Matrix &velocity_x, const Matrix &velocity_y, c
    int N = grid.size(); // number of grid points
    Matrix new_grid = grid; // create copy of current grid

    // Parallelize the grid update
    #pragma omp parallel for collapse(2)
    for (int i = 1; i < (N - 1); ++i) {
        for (int j = 1; j < (N - 1); ++j) { // iterate over the interior grid points

            // compute advection term using upwind differencing
            double advection_x;
            if (velocity_x[i][j] > 0) {
                advection_x = velocity_x[i][j] * (grid[i][j] - grid[i][j-1]) / dx;
            } else {
                advection_x = velocity_x[i][j] * (grid[i][j+1] - grid[i][j]) / dx;
            }
        }
    }
```

# Parallelization with OpenMP

## Parallelizing AMR Error Calculation:

- Since each grid cell or quadtree leaf is processed independently during AMR, you can use OpenMP to parallelize the error calculation and refinement/coarsening

#pragma omp parallel for collapse(2)

- Speeding up traversal in the nested loops

#pragma omp critical

- For thread-safe updates to the quadtree

```
void apply_amr(Node* root, const Matrix &grid, double dx, double dy, double refine_threshold, double coarsen_threshold) {
    #pragma omp parallel for collapse(2)
    for (int i = 1; i < grid.size() - 1; ++i) { // loop through interior cell
        for (int j = 1; j < grid[0].size() - 1; ++j) {
            // calculate the error at the current grid cell
            double error = calculate_error(grid, i, j, dx, dy);

            #pragma omp critical // lock to ensure thread-safety
            {
                // refine the cell if the error exceeds the refine threshold
                if (error > refine_threshold) {
                    refine_grid(root);
                }
                // coarsen the cell if the error is below the coarsen threshold
                else if (error < coarsen_threshold) {
                    coarsen_grid(root);
                }
            }
        }
    }
}
```

# Parallelization with OpenMP

## Parallelizing Max Velocity Calculation:

#pragma omp reduction(max : max\_velocity)

- The reduction clause ensures that OpenMP combines the local maxima from all threads into max\_velocity after the loop

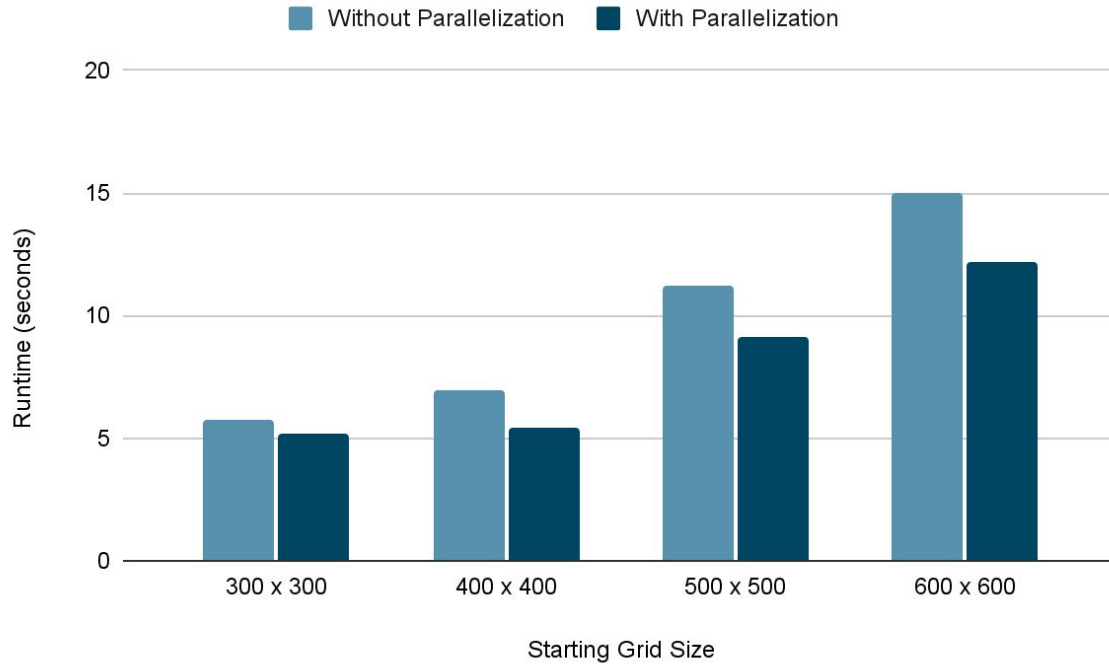
```
// calculate stable time step
double max_velocity = 0.0;
#pragma omp parallel for reduction(max:max_velocity)
for (int i = 0; i < N; ++i) {
    for (int j = 0; j < N; ++j) {
        max_velocity = max(max_velocity, max(fabs(velocity_x[i][j]), fabs(velocity_y[i][j])));
    }
}
```

# Scalability Results


Starting Grid Size	Without Parallelization	With Parallelization (OpenMP)
300 x 300	5.75946 s	5.21159 s
400 x 400	6.92825 s	5.40180 s
500 x 500	11.2233 s	10.0942 s
600 x 600	15.0424 s	13.2068 s

Note: I used four threads

# Scalability Results



# Final Results





The background is a dark navy blue. In the top-left corner, there are several overlapping, rounded, organic shapes in a slightly lighter shade of blue, creating a layered effect. A thin, elegant gold line curves from the bottom-left towards the bottom-right, entering the frame from the left and exiting on the right.

**Questions?**